

Librerías

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import f1_score, mean_squared_error, confusion_matrix, ConfusionMatrixDisplay
from sklearn.impute import KNNImputer
from sklearn.neural_network import MLPClassifier
from sklearn.preprocessing import LabelEncoder, StandardScaler, OneHotEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from imblearn.over_sampling import RandomOverSampler
```

Obtención Datos

```
In [ ]: oTrainData = pd.read_csv("train.csv")
oTestData = pd.read_csv("test.csv")
```

Analisis de datos

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt

# Cargar el conjunto de datos
oTrainData = pd.read_csv('/content/train.csv')

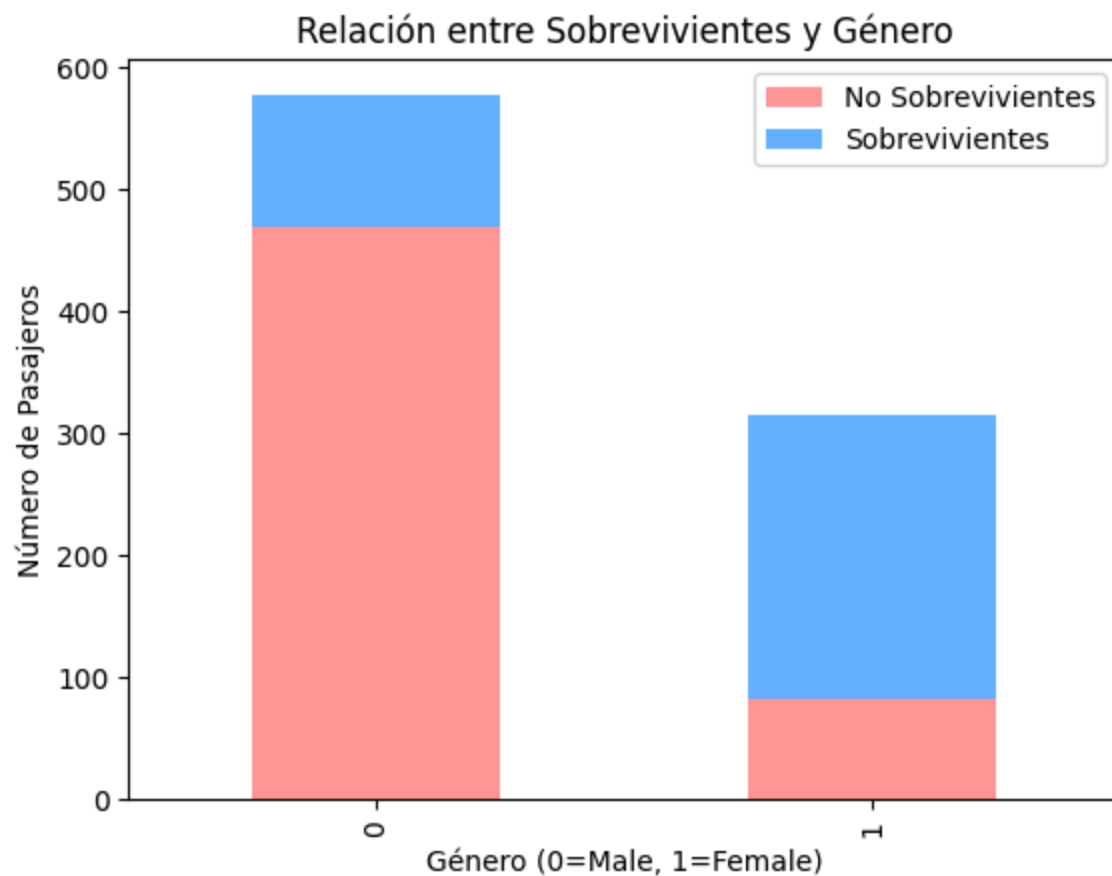
# Mapear 'female' a 1 y 'male' a 0 en la columna 'Sex'
oTrainData['Sex'] = oTrainData['Sex'].map({'female': 1, 'male': 0})

# Rellenar los valores faltantes en la columna 'Age' con la edad media
mean_age = oTrainData['Age'].mean()
oTrainData['Age'].fillna(mean_age, inplace=True)

# Verificar que la columna 'Survived' exista en el DataFrame
if 'Survived' not in oTrainData.columns:
    raise ValueError("La columna 'Survived' no se encuentra en el DataFrame.")

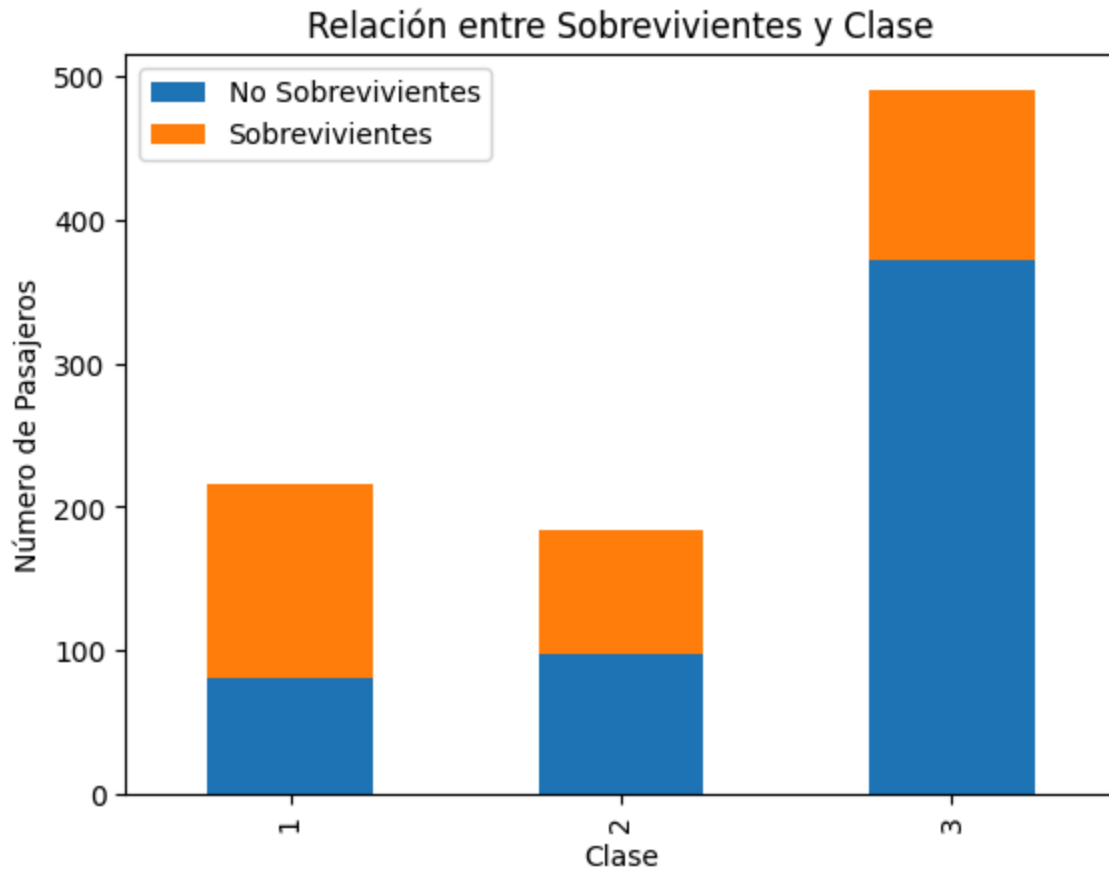
# Agrupar por género y sobrevivencia para contar la cantidad de sobrevivientes y no sobrevivientes
survival_by_gender = oTrainData.groupby(['Sex', 'Survived']).size().unstack()

# Crear la gráfica de barras
survival_by_gender.plot(kind='bar', stacked=True, color=['#ff9999', '#66b3ff'])
plt.xlabel('Género (0=Male, 1=Female)')
plt.ylabel('Número de Pasajeros')
plt.title('Relación entre Sobrevivientes y Género')
plt.legend(['No Sobrevivientes', 'Sobrevivientes'])
plt.show()
```



```
In [ ]: import matplotlib.pyplot as plt
survival_by_class = oTrainData.groupby(['Pclass', 'Survived']).size().unstack()

# Crear la gráfica de barras
survival_by_class.plot(kind='bar', stacked=True)
plt.xlabel('Clase')
plt.ylabel('Número de Pasajeros')
plt.title('Relación entre Sobrevivientes y Clase')
plt.legend(['No Sobrevivientes', 'Sobrevivientes'])
plt.show()
```



```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt

survival_by_class_gender = oTrainData.groupby(['Pclass', 'Sex', 'Survived']).size().unstack()

# Calcular el total de personas en cada clase social y género
total_by_class_gender = survival_by_class_gender.sum(axis=1)

# Agregar el total al DataFrame survival_by_class_gender
survival_by_class_gender['Total'] = total_by_class_gender

# Mostrar la cantidad de personas que sobrevivieron, murieron y el total por clase social y género
print("Cantidad de personas que sobrevivieron, murieron y el total por clase social y género")
print(survival_by_class_gender)

# Calcula el índice de supervivencia por clase social y género
survival_rate_by_class_gender = survival_by_class_gender[1] / survival_by_class_gender['Total']

# Redondear los índices de supervivencia a enteros (porcentaje)
survival_rate_by_class_gender_percent = (survival_rate_by_class_gender * 100).round().astype(int)

# Mostrar el índice de supervivencia por clase social y género en porcentaje
print("\nÍndice de supervivencia por clase social y género (redondeado):")
print(survival_rate_by_class_gender_percent)

# Graficar el índice de supervivencia por clase social y género
ax = survival_rate_by_class_gender_percent.unstack().plot(kind='bar', color=['skyblue', 'red'])

# Añadir las anotaciones con el índice de supervivencia redondeado
for container in ax.containers:
    ax.bar_label(container, label_type='edge')
```

```
# Ajustar etiquetas y título
plt.xlabel('Clase social y Género')
plt.ylabel('Índice de supervivencia (%)')
plt.title('Índice de supervivencia por clase social y género con totales')
plt.xticks(rotation=45)
plt.legend(['Hombres', 'Mujeres'])

# Mostrar la gráfica
plt.show()
```

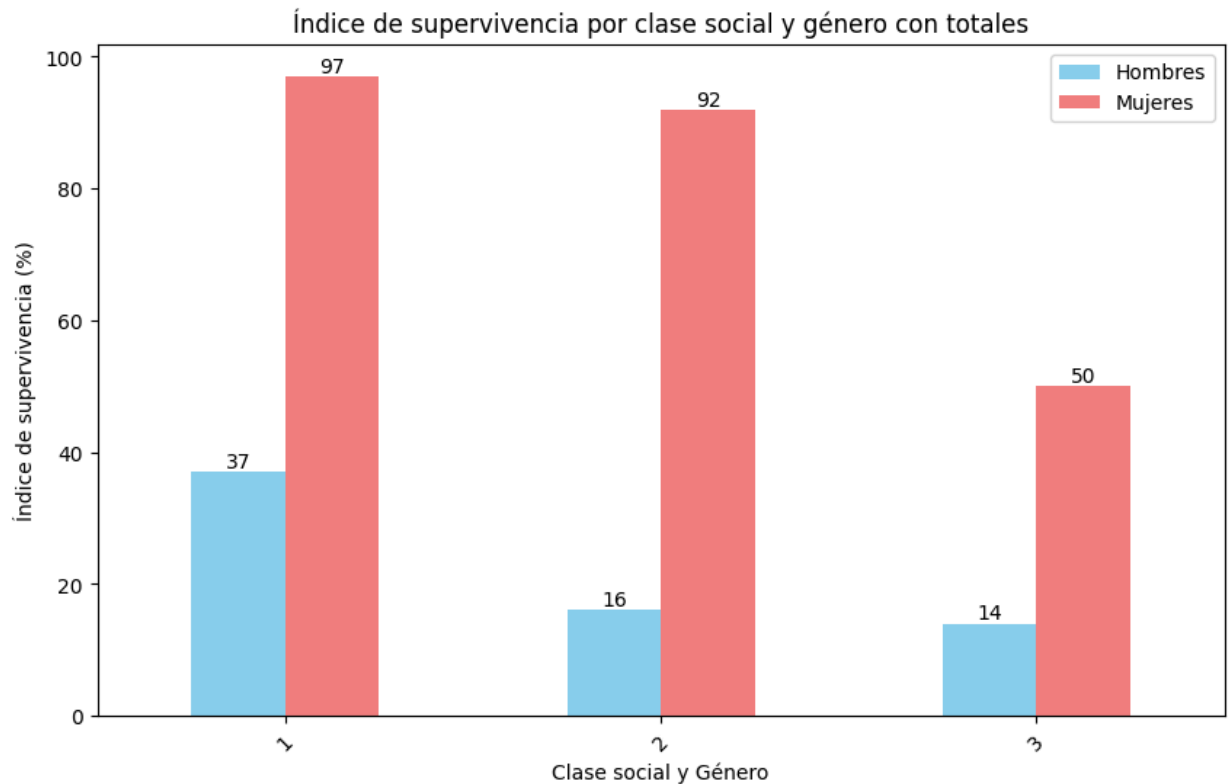
Cantidad de personas que sobrevivieron, murieron y el total por clase social y género:

Survived	0	1	Total
Pclass Sex			
1 0	77	45	122
1 1	3	91	94
2 0	91	17	108
2 1	6	70	76
3 0	300	47	347
3 1	72	72	144

Índice de supervivencia por clase social y género (redondeado):

Pclass	Sex	
1	0	37
1	1	97
2	0	16
2	1	92
3	0	14
3	1	50

dtype: int64



```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt
# Define las edades
```

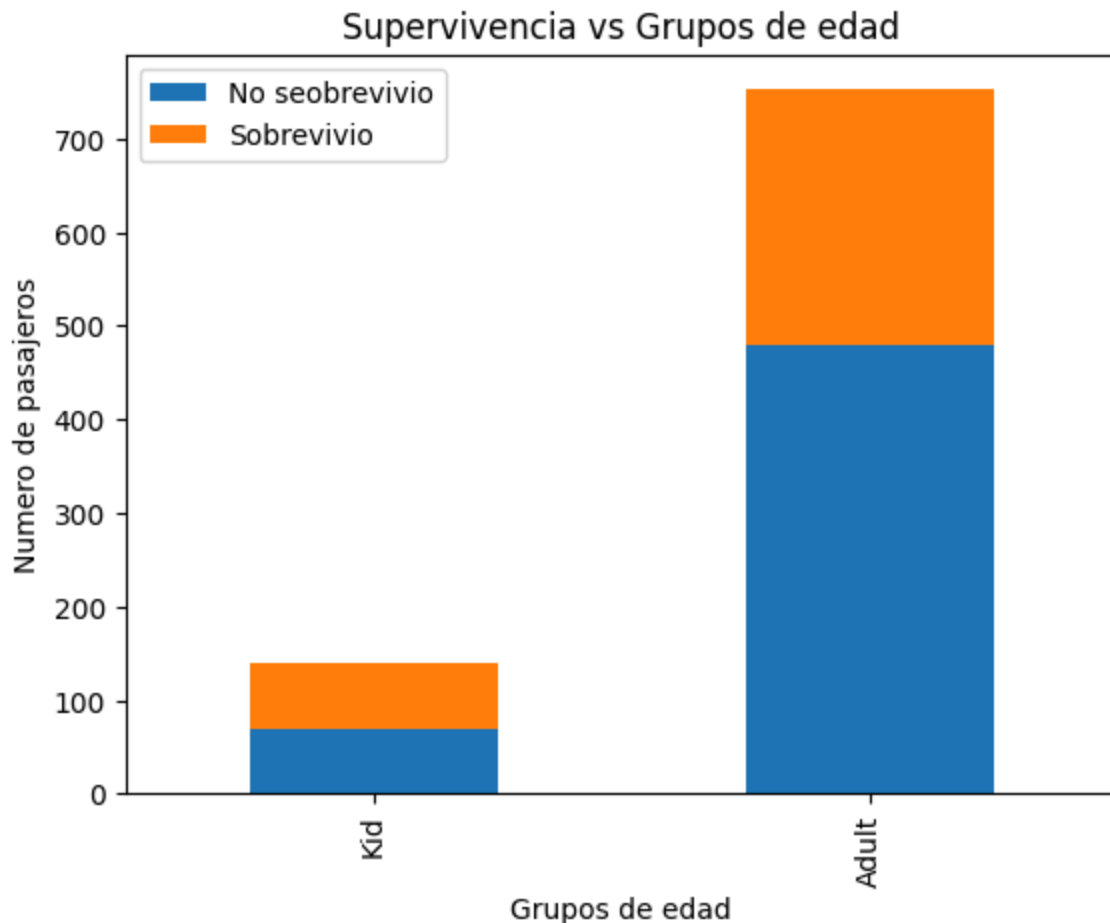
```
oTrainData['AgeGroup'] = pd.cut(oTrainData['Age'], bins=[0, 18, 100], labels=['Kid', 'Adult'])

survival_by_age_group = oTrainData.groupby(['AgeGroup', 'Survived']).size().unstack()

survival_by_age_group.plot(kind='bar', stacked=True)
plt.xlabel('Grupos de edad')
plt.ylabel('Numero de pasajeros')
plt.title('Supervivencia vs Grupos de edad')
plt.legend(['No seobrevivio', 'Sobrevivio'])
plt.show()
```

<ipython-input-13-af080d706e4a>:6: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
survival_by_age_group = oTrainData.groupby(['AgeGroup', 'Survived']).size().unstack()
```



```
In [ ]: import pandas as pd

oTrainData = pd.read_csv('/content/train.csv')

# Limpiar datos: eliminar filas con edad faltante
oCleanData = oTrainData[~oTrainData['Age'].isnull()].copy()
oCleanData.index = range(len(oCleanData))

# Definir grupos de edad
oCleanData.loc[:, 'AgeGroup'] = pd.cut(oCleanData['Age'], bins=range(0, 101, 5), labels=['Kid', 'Adult'])

# Mapear 'female' a 1 y 'male' a 0 en la columna 'Sex'
```

```
oCleanData.loc[:, 'Sex'] = oCleanData['Sex'].map({'female': 1, 'male': 0})

# Agrupar por grupo de edad, género y supervivencia
survival_by_age_gender = oCleanData.groupby(['AgeGroup', 'Sex', 'Survived'], observed=
# Calcular el total de personas por grupo de edad
total_by_age_group = survival_by_age_gender.sum(axis=1)

# Crear una tabla que combine total, fallecidos y sobrevivientes desglosado por género
summary_by_age_gender = pd.DataFrame({
    'Total': total_by_age_group,
    'Hombres Fallecidos': survival_by_age_gender[(0, 0)],
    'Hombres Sobrevivientes': survival_by_age_gender[(0, 1)],
    'Mujeres Fallecidas': survival_by_age_gender[(1, 0)],
    'Mujeres Sobrevivientes': survival_by_age_gender[(1, 1)]
})

# Mostrar la tabla con el resumen
print("Cantidad total, fallecidos y sobrevivientes por grupo de edad (desglosado por g
print(summary_by_age_gender.to_string())

# Calcular el índice de supervivencia por grupo de edad
survival_rate_by_age_group = survival_by_age_gender[(1, 1)] / total_by_age_group

# Mostrar el índice de supervivencia por grupo de edad
print("\nÍndice de supervivencia por grupo de edad:")
print(survival_rate_by_age_group)


# Graficar el índice de supervivencia por grupo de edad
survival_rate_by_age_group.plot(kind='bar', color='skyblue')
plt.xlabel('Grupo de edad')
plt.ylabel('Índice de supervivencia')
plt.title('Índice de supervivencia por grupo de edad')
plt.show()
```

Cantidad total, fallecidos y sobrevivientes por grupo de edad (desglosado por género):

	Total	Hombres Fallecidos	Hombres Sobrevivientes	Mujeres Fallecidas	Mujeres Sobrevivientes
AgeGroup					
0-4	44	8	15	5	16
5-9	20	6	4	7	3
10-14	19	5	2	3	9
15-19	96	54	8	9	25
20-24	122	69	8	11	34
25-29	108	57	15	9	27
30-34	88	43	15	4	26
35-39	67	34	8	5	20
40-44	47	23	6	7	11
45-49	39	20	6	3	10
50-54	24	14	2	0	8
55-59	18	10	2	1	5
60-64	14	10	1	0	3
65-69	3	3	0	0	0
70-74	4	4	0	0	0
75-79	1	0	1	0	0

Índice de supervivencia por grupo de edad:


AgeGroup	
0-4	0.363636
5-9	0.150000
10-14	0.473684
15-19	0.260417
20-24	0.278689
25-29	0.250000
30-34	0.295455
35-39	0.298507
40-44	0.234043
45-49	0.256410
50-54	0.333333
55-59	0.277778
60-64	0.214286
65-69	0.000000
70-74	0.000000
75-79	0.000000
dtype:	float64



```
In [ ]: import matplotlib.pyplot as plt
# Create a new column for family size
oTrainData['FamilySize'] = oTrainData['SibSp'] + oTrainData['Parch'] + 1

# Group by family size and survival
survival_by_family_size = oTrainData.groupby(['FamilySize', 'Survived']).size().unstack()


# Plot bar chart
survival_by_family_size.plot(kind='bar', stacked=True)
plt.xlabel('Family Size')
plt.ylabel('Number of Passengers')
plt.title('Survival vs Family Size')
plt.legend(['Did Not Survive', 'Survived'])
plt.show()
```

```
In [ ]: # Supervivencia vs Grupos de edad
survival_by_parch = oCleanData.groupby(['Parch', 'Survived']).size().unstack()

# Mostrar el índice de supervivencia por grupo de edad
total_by_parch = survival_by_parch.sum(axis=1)
survival_by_parch_index = survival_by_parch.div(total_by_parch, axis=0) * 100


# Graficar el índice de supervivencia por grupo de edad
survival_by_parch.plot(kind='bar', stacked=True, color=['#ff9999', '#66b3ff'])
plt.xlabel('Número de padres o hijos')
plt.ylabel('Índice de supervivencia')
plt.title('Índice de supervivencia por parch')
plt.show()
```



```
In [ ]: # Supervivencia vs Grupos de edad
survival_by_sibsp = oCleanData.groupby(['SibSp', 'Survived']).size().unstack()

# Mostrar el índice de supervivencia por grupo de edad
total_by_sibsp = survival_by_sibsp.sum(axis=1)
survival_by_sibsp_index = survival_by_sibsp.div(total_by_parch, axis=0) * 100

# Graficar el índice de supervivencia por grupo de edad
survival_by_sibsp.plot(kind='bar', stacked=True, color=['#ff9999', '#66b3ff'])
plt.xlabel('Número de hermanos o esposos')
plt.ylabel('Índice de hermanos o esposos')
plt.title('Índice de supervivencia por número de hermanos o esposos')
plt.show()
```



```
In [ ]: bins = [0, 18] + list(range(19, 101, 5))

# Create labels for each age group
labels = ['Niño'] + [f'{i}-{i+4}' for i in range(19, 100, 5)]

# Categorize the ages into the defined bins
oTrainData['AgeClass'] = pd.cut(oTrainData['Age'], bins=bins, labels=labels, right=False)

# Calculate the counts and percentages
age_class_counts = oTrainData.groupby(['AgeClass', 'Pclass'])['PassengerId'].count().unstack()
age_class_percentages = age_class_counts.div(age_class_counts.sum(axis=1), axis=0) * 100

# Plot the data
ax = age_class_percentages.plot(kind='bar', stacked=True, figsize=(10, 6))


# Add labels to the bars
for container in ax.containers:
    ax.bar_label(container, label_type='center', fmt='%.0f%%')

# Configure the labels and title
plt.xlabel('Grupo de Edad')
plt.ylabel('Porcentaje de Pasajeros')
plt.title('Distribución de Clases por Grupo de Edad')
plt.xticks(rotation=45, ha='right')
plt.legend(title='Clase', loc='upper right')

# Show the plot
plt.show()
```

```
<ipython-input-18-d2412e5a75d6>:10: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.
```

```
age_class_counts = oTrainData.groupby(['AgeClass', 'Pclass'])['PassengerId'].count().unstack()
```




```
In [ ]: oTrainData['FamilySize'] = oTrainData['SibSp'] + oTrainData['Parch'] + 1

survival_by_family_gender = oTrainData.groupby(['FamilySize', 'Sex', 'Survived']).size

total_by_family_gender = survival_by_family_gender.sum(axis=1)

survival_rate_by_family_gender = survival_by_family_gender[1] / total_by_family_gender

survival_rate_by_family_gender.unstack().plot(kind='bar', figsize=(10, 6))
plt.xlabel('Tamaño de la Familia')
plt.ylabel('Índice de Supervivencia')
plt.title('Índice de Supervivencia en Base al Tamaño de la Familia y Género')
plt.xticks(rotation=0)
plt.legend(['Masculino', 'Femenino'])
plt.show()
```



```
In [ ]: oTrainData_adults = oTrainData[oTrainData['Age'] >= 18]

oTrainData_adults['FamilySize'] = oTrainData_adults['SibSp'] + oTrainData_adults['ParCh']

survival_by_family_gender = oTrainData_adults.groupby(['FamilySize', 'Sex', 'Survived'])
total_by_family_gender = survival_by_family_gender.sum(axis=1)
survival_rate_by_family_gender = survival_by_family_gender[1] / total_by_family_gender

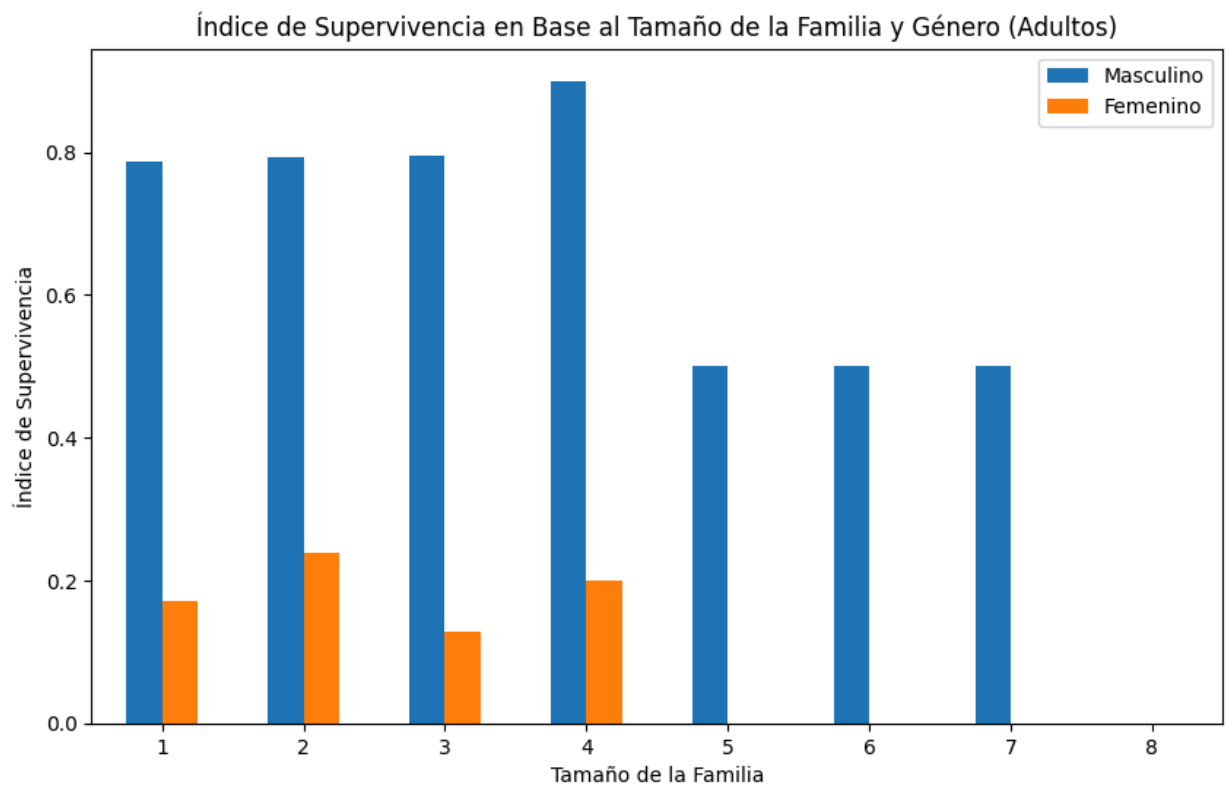
survival_rate_by_family_gender.unstack().plot(kind='bar', figsize=(10, 6))
plt.xlabel('Tamaño de la Familia')
plt.ylabel('Índice de Supervivencia')
plt.title('Índice de Supervivencia en Base al Tamaño de la Familia y Género (Adultos)')
plt.xticks(rotation=0)
plt.legend(['Masculino', 'Femenino'])
plt.show()
```

<ipython-input-20-d8834f234bc7>:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
oTrainData_adults['FamilySize'] = oTrainData_adults['SibSp'] + oTrainData_adults['ParCh'] + 1
```



```
In [ ]: import matplotlib.pyplot as plt
# supervivencia en base al puerto de embarque
survival_by_embarked = oTrainData.groupby(['Embarked', 'Survived']).size().unstack(fill_value=0)
total_by_embarked = survival_by_embarked.sum(axis=1)
survival_rate_by_embarked = survival_by_embarked[1] / total_by_embarked

survival_rate_by_embarked.plot(kind='bar', figsize=(10, 6))
plt.xlabel('Puerto de Embarque')
plt.ylabel('Índice de Supervivencia')
plt.title('Índice de Supervivencia en Base al Puerto de Embarque')
plt.xticks(rotation=0)
plt.show()
```



```
In [ ]: import matplotlib.pyplot as plt
# Agrupar por puerto de embarque, género y supervivencia
survival_by_embarked_gender = oTrainData.groupby(['Embarked', 'Sex', 'Survived']).size

# Calcular el total de personas por puerto de embarque y género
total_by_embarked_gender = survival_by_embarked_gender.sum(axis=1)

# Calcular el índice de supervivencia por puerto de embarque y género
survival_rate_by_embarked_gender = survival_by_embarked_gender[1] / total_by_embarked_gender

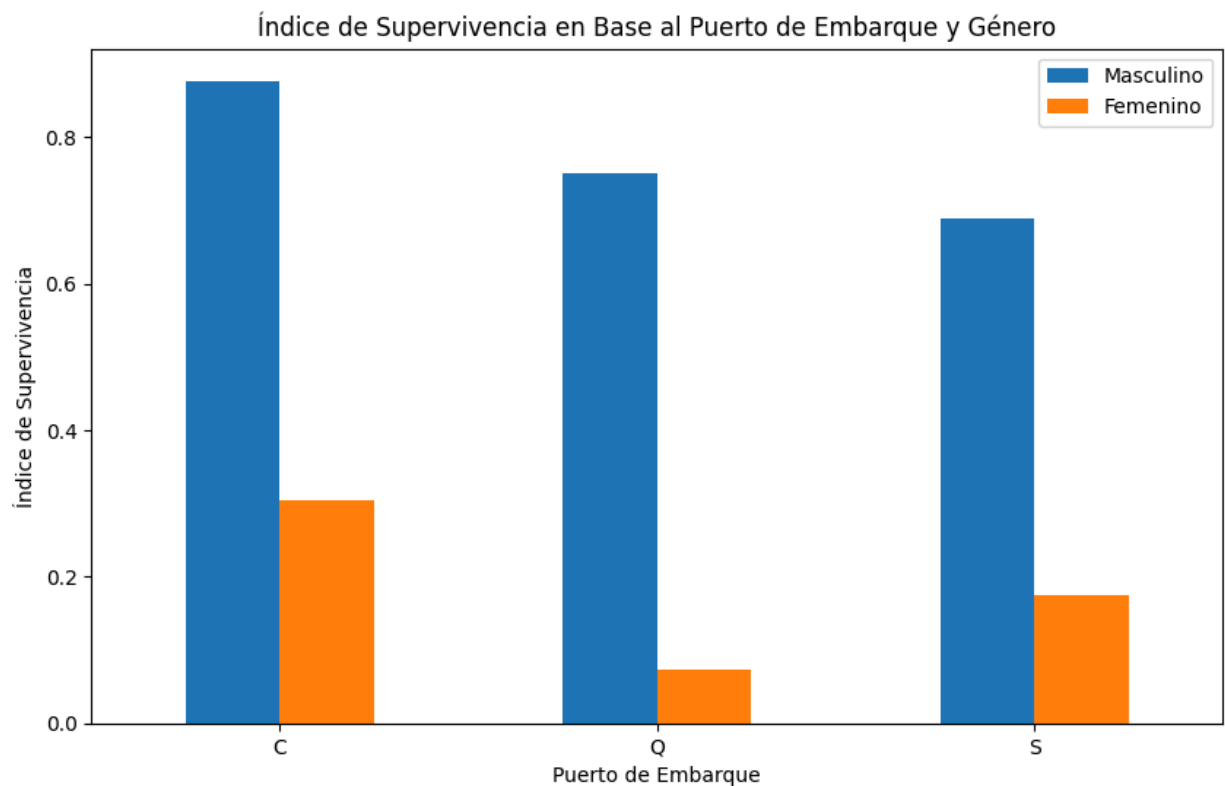
# Mostrar el índice de supervivencia por puerto de embarque y género
print("\nÍndice de supervivencia por puerto de embarque y género:")
print(survival_rate_by_embarked_gender)

# Graficar el índice de supervivencia por puerto de embarque y género
survival_rate_by_embarked_gender.unstack().plot(kind='bar', figsize=(10, 6))
plt.xlabel('Puerto de Embarque')
plt.ylabel('Índice de Supervivencia')
plt.title('Índice de Supervivencia en Base al Puerto de Embarque y Género')
plt.xticks(rotation=0)
plt.legend(['Masculino', 'Femenino'])
plt.show()
```

Índice de supervivencia por puerto de embarque y género:

Embarked	Sex	
C	female	0.876712
	male	0.305263
Q	female	0.750000
	male	0.073171
S	female	0.689655
	male	0.174603

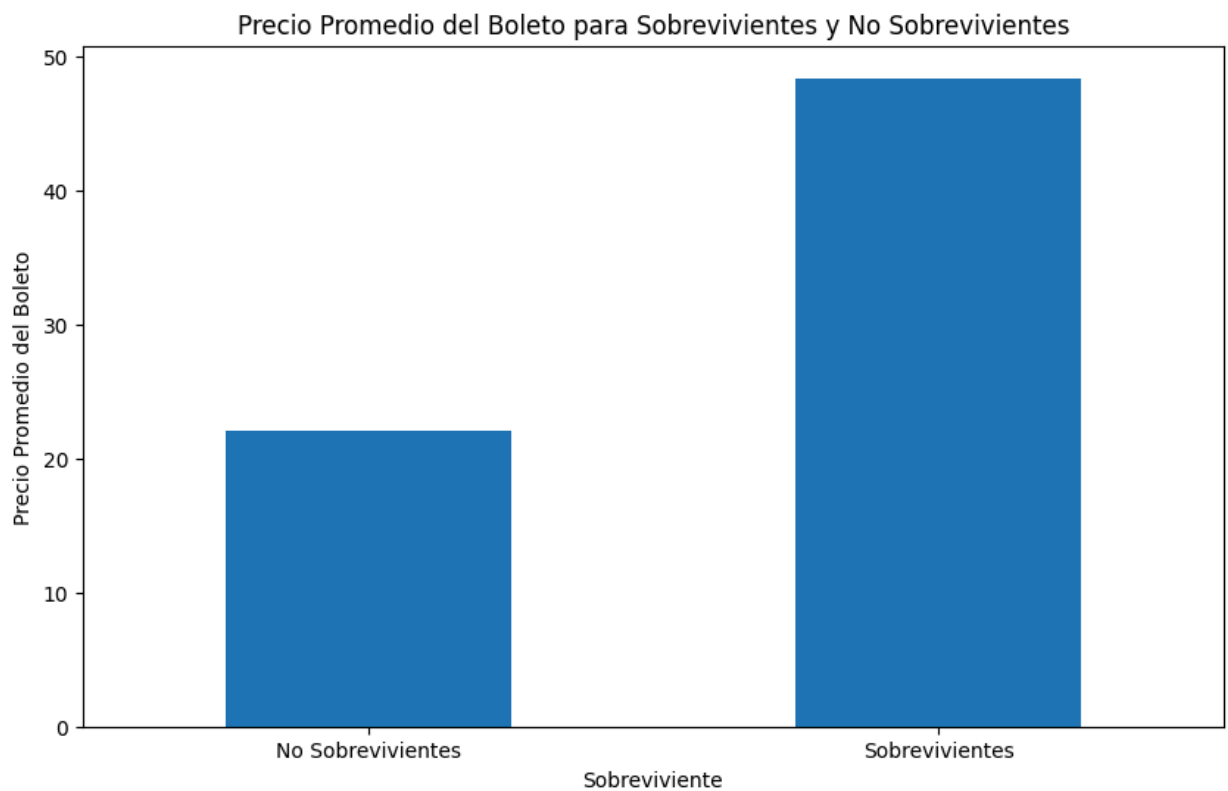
dtype: float64



```
In [ ]: import matplotlib.pyplot as plt

# supervivencia en base al precio del boleto
survival_by_fare = oTrainData.groupby(['Survived'])['Fare'].mean()


survival_by_fare.plot(kind='bar', figsize=(10, 6))
plt.xlabel('Sobreviviente')
plt.ylabel('Precio Promedio del Boleto')
plt.title('Precio Promedio del Boleto para Sobrevivientes y No Sobrevivientes')
plt.xticks([0, 1], ['No Sobrevivientes', 'Sobrevivientes'], rotation=0)
plt.show()
```

```
In [ ]: import matplotlib.pyplot as plt

# supervivencia en base a si el pasajero tenía cabina asignada
oTrainData['HasCabin'] = oTrainData['Cabin'].notnull().astype(int)
survival_by_cabin = oTrainData.groupby(['HasCabin', 'Survived']).size().unstack(fill_v
total_by_cabin = survival_by_cabin.sum(axis=1)
survival_rate_by_cabin = survival_by_cabin[1] / total_by_cabin

survival_rate_by_cabin.plot(kind='bar', figsize=(10, 6))
plt.xlabel('Cabina Asignada')
plt.ylabel('Índice de Supervivencia')
plt.title('Índice de Supervivencia en Base a si el Pasajero Tenía Cabina Asignada')
plt.xticks([0, 1], ['Sin Cabina', 'Con Cabina'], rotation=0)
plt.show()
```



Limpieza de Datos

```
In [ ]: # 1 niños, 2 adultos, 3 adultos mayores
def label_age(age):
    if 0 <= age <= 10: return 1
    elif 11 <= age <= 20: return 2
    elif 21 <= age <= 35: return 3
    elif 36 <= age <= 50: return 4
    elif 51 <= age <= 65: return 5
    elif 66 <= age <= 80: return 6
    else: return 0

def clean_data(df, age_neighbors = False, nbs_number = 4, group_age = True, scale_age=
    scaler = StandardScaler()
    hot_encoder = OneHotEncoder()

    data_copy = df.copy().drop(columns = ["PassengerId", "Name", "Fare", "Cabin", "Ticket"])
    data_copy["Sex"] = data_copy["Sex"].map({"male": 0, "female": 1})
    data_copy["Embarked"].fillna(data_copy["Embarked"].mode()[0], inplace = True)

    encoded_embarked = hot_encoder.fit_transform(data_copy[["Embarked"]]).toarray()
    encoded_embarked = pd.DataFrame(encoded_embarked, columns = ["C", "Q", "S"])
    data_copy = pd.concat([data_copy, encoded_embarked], axis = 1)
    data_copy.drop(columns = ["Embarked"], inplace = True)

    #data_copy["Fare"].fillna(data_copy["Fare"].median(), inplace = True)
    #data_copy[["Fare"]] = scaler.fit_transform(data_copy[["Fare"]])

    if age_neighbors:
        data_for_imputation = data_copy[["Age", "Pclass", "Sex", "SibSp", "Parch", "C", "Q", "S"]]
        imputer = KNNImputer(n_neighbors = nbs_number)
        imputed_data = imputer.fit_transform(data_for_imputation)
```

```

imputed_df = pd.DataFrame(imputed_data, columns = data_for_imputation.columns)
data_copy["Age"] = imputed_df["Age"]

else:
    data_copy["Age"].fillna(data_copy["Age"].median(), inplace = True)

if group_age:
    data_copy["Age"] = data_copy["Age"].apply(label_age)

if scale_age:
    data_copy[["Age"]] = scaler.fit_transform(data_copy[["Age"]])

return data_copy

def oversample(x, y):
    ros = RandomOverSampler()
    x, y = ros.fit_resample(x, y)
    return x, y

```

```

In [ ]: #oTrainData = pd.read_csv("train.csv")
        #oTestData = pd.read_csv("test.csv")

oCleanTrainData = clean_data(oTrainData, age_neighbors=True, group_age=False, scale_age=True)
#oCleanTrainData.head()
oCleanTrainData, oCleanTestData = train_test_split(oCleanTrainData, test_size=0.2)

oCleanTrainData, oCleanValidationData = train_test_split(oCleanTrainData, test_size=0.2)

print("Train Length: ", len(oCleanTrainData))
print("Validation Length: ", len(oCleanValidationData))
print("Test Length: ", len(oCleanTestData))

Train Length: 356
Validation Length: 356
Test Length: 179

```

```
In [ ]: oCleanTrainData.head()
```

```
Out[ ]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	C	Q	S
721	0	3	0	-0.891171	1	0	0.0	0.0	1.0
354	0	3	0	-0.386324	0	0	1.0	0.0	0.0
135	0	2	0	-0.450577	0	0	1.0	0.0	0.0
694	0	1	0	2.266420	0	0	0.0	0.0	1.0
466	0	2	0	0.650908	0	0	0.0	0.0	1.0

Modelo Clasificacion

```

In [ ]: oFScore = []
        train_sizes = np.arange(0.2, 1.0, 0.1) # 40% a 90%, incrementando de 10 en 10

```

```

best_results = []
best_model = None
best_solver = None
oSolverList = ["lbfgs", "liblinear", "newton-cg", "newton-cholesky", "sag", "saga"]
results = []

for sSolver in oSolverList:
    for train_size in train_sizes:

        best_seed, best_score, best_precision = None, float("-inf"), float("-inf")
        best_score = 0

        seeds = [[np.random.randint(0, 2**32 - 1), np.random.randint(0, 2**32 - 1)] for _
        for seed in seeds:

            #Se asigna el tamaño del Sample con el que se va a entrenar
            iSampleSize = int( len(oCleanTrainData) * train_size )

            #Se separa el subset de entrenamiento por el tamaño que se asigno arriba
            oTrainSubset = oCleanTrainData.sample(n=iSampleSize)
            x, y = oversample(oTrainSubset[["Age", "Pclass", "Sex", "SibSp", "Parch", "C", "Q"]])

            #se entrena el modelo en base a los hiperparametros de Semilla y Solver
            oModelTemp = LogisticRegression(random_state=seed[1], solver=sSolver, max_iter=1000)
            #oModelTemp.fit(oTrainSubset[["Age", "Pclass", "Sex", "SibSp", "Parch", "C", "Q"]])
            oModelTemp.fit(x,y)

            #Se generan las predicciones en base a los datos de Entrenamiento con nuestro modelo
            oTrainingPredict = oModelTemp.predict(oTrainSubset[["Age", "Pclass", "Sex", "SibSp", "Parch", "C", "Q"]])
            fTrainingFScore = f1_score(oTrainSubset["Survived"], oTrainingPredict)
            fTrainingLogLoss = log_loss(oTrainSubset["Survived"], oTrainingPredict)

            #Se generan las predicciones en base a los datos de Validacion con modelo ya entrenado
            oValidationPredict = oModelTemp.predict(oCleanValidationData[["Age", "Pclass", "Sex", "SibSp", "Parch", "C", "Q"]])
            fValidationFScore = f1_score(oCleanValidationData["Survived"], oValidationPredict)
            fValidationLogLoss = log_loss(oCleanValidationData["Survived"], oValidationPredict)

            #Por si acaso guardamos el error cuadrático
            oFScore.append(fValidationFScore)

            #Si el error cuadrático, de las predicciones de Validacion, es menor al anterior
            if fValidationFScore > best_score:
                best_score = fValidationFScore
                best_seed = seed
                best_precision = precision_score(oCleanValidationData["Survived"], oValidationPredict)
                best_model = oModelTemp # Guardar el mejor modelo entrenado
                best_solver = sSolver # Guarda con que Solver tuvo el mejor resultado
                best_results.append({
                    "train_size": train_size,
                    "best_seed_split": best_seed[0],
                    "best_seed_train": best_seed[1],
                    "best_precision": best_precision,
                    "best_score": best_score,
                    "best_solver": best_solver,
                })

            #Guardamos todos los resultados
            results.append((train_size, sSolver, fTrainingFScore, fTrainingLogLoss, fValidationFScore, fValidationLogLoss))

```

```
# Convertir resultados a DataFrame y exportar a CSV
logistic_data = pd.DataFrame(results, columns = ["train_size", "solver", "fTrainingFScore"])
best_results_df = pd.DataFrame(best_results)
best_results_df.to_csv("best_tree_results.csv", index=False)

# Mostrar los mejores resultados
print(best_results_df)
```

	train_size	best_seed_split	best_seed_train	best_precision	best_score	\
0	0.2	1723463638	177801891	0.703947	0.735395	
1	0.2	2083933156	3076909759	0.746479	0.754448	
2	0.2	3687173766	371767915	0.773050	0.778571	
3	0.2	2403670650	1809800909	0.778571	0.781362	
4	0.3	4179813320	640835607	0.762238	0.773050	
..	
227	0.8	2191697976	2595254971	0.750000	0.763251	
228	0.8	2417631648	924193493	0.758621	0.774648	
229	0.8	553152310	20805895	0.767606	0.775801	
230	0.9	2262574612	4044196783	0.756944	0.770318	
231	0.9	3635354591	3505707082	0.758621	0.774648	

	best_solver
0	lbfgs
1	lbfgs
2	lbfgs
3	lbfgs
4	lbfgs
..	...
227	saga
228	saga
229	saga
230	saga
231	saga


[232 rows x 6 columns]

```
In [ ]: best_index = best_results_df["best_score"].idxmax()
best_model = best_model # Utiliza el modelo del mejor F1 score global
test_predicciones = best_model.predict(oCleanTestData[["Age", "Pclass", "Sex", "SibSp", "Survived"]])


fTestFScore = f1_score(oCleanTestData["Survived"], test_predicciones)
print("F1 Score: ", fTestFScore)
print("Log Loss: ", log_loss(oCleanTestData["Survived"], test_predicciones))
print("Precision: ", precision_score(oCleanTestData["Survived"], test_predicciones))

cm = confusion_matrix(oCleanTestData["Survived"], test_predicciones)
disp = ConfusionMatrixDisplay(confusion_matrix = cm, display_labels = best_model.classes_)
disp.plot(cmap = plt.cm.Blues)
plt.show()
```

F1 Score: 0.7272727272727272
 Log Loss: 7.249002916247025
 Precision: 0.7619047619047619



```
In [ ]: plt.figure(figsize=(10, 6))
plt.scatter(logistic_data["train_size"] , logistic_data["fTrainingFScore"], label='Train')
plt.scatter(logistic_data["train_size"] , logistic_data["fValidationFScore"], label='Validation')
plt.axhline(y=fTestFScore, color='red', linestyle='--', label='Test Error (Base)')
plt.xlabel('Tamaño de muestra')
plt.ylabel('Error cuadrático medio')
plt.title('Evolución del error promedio de entrenamiento y validación')
plt.legend()
plt.show()
```



```
In [ ]: logistic_data = logistic_data.sort_values(by = "fValidationFScore", ascending = True)
logistic_data.head()
```

```
Out [ ]:
```

	train_size	solver	fTrainingFScore	fTrainingLogLoss	fValidationFScore	fValidationLogLoss
1694	0.2	newton-cg	0.777778	6.091885	0.611321	10.428360
1662	0.2	newton-cg	0.727273	6.091885	0.629032	9.314652
158	0.3	lbfgs	0.729730	6.800689	0.635659	9.517144
3244	0.2	sag	0.763636	6.599542	0.636735	9.010913
3357	0.3	sag	0.753623	5.780586	0.637795	9.314652

```
In [ ]: oTestData = pd.read_csv('/content/test.csv')

oSubmissionTestData = clean_data(oTestData)
best_index = best_results_df["best_score"].idxmax()
best_model = best_model # Utiliza el modelo del mejor F1 score global
test_predicciones = best_model.predict(oSubmissionTestData[["Age", "Pclass", "Sex", "S

# Crear el archivo de salida con las predicciones
submit_data = pd.DataFrame(oTestData, columns=["PassengerId"])
submit_data["Survived"] = test_predicciones
submit_data.to_csv("prediction.csv", index=False)
print("Predicciones guardadas en 'prediction.csv'")
```

Predicciones guardadas en 'prediction.csv'

Modelo Redes Neuronales

```
In [ ]: def train_nn_model(train_x, train_y, layers, alpha, solver, max_iter, learning_rate_val):
        nn = MLPClassifier(activation=activation_value, hidden_layer_sizes=layers, max_iter=
            solver=solver, alpha=alpha, learning_rate="adaptive", learning_rate_val=learning_rate_val)
        nn.fit(train_x, train_y)
        return nn
```

```
In [ ]: layers = [(i, j, k, 1) for i in range(2, 7) for j in range(2, 5) for k in range(2, 4)]
train_sizes = np.arange(0.2, 1.0, 0.1) # 40% a 90%, incrementando de 10 en 10
solvers = ["adam", "sgd", "lbfgs"]

oNNScore = []
results = []
best_results = []
best_nn_model = None
best_score, best_log_loss = float("-inf"), float("inf")
MAX_ITER = 2000

for layer in layers:
    for train_size in train_sizes:
        for solver in solvers:
            train_x, _, train_y, _ = train_test_split(
                oCleanTrainData.drop(columns=["Survived"]), oCleanTrainData["Survived"], train_size=train_size,
                train_x, train_y = oversample(train_x, train_y)

            nn_model = MLPClassifier(hidden_layer_sizes=layer, max_iter=MAX_ITER, solver=solver)
            nn_model.fit(train_x, train_y)

            # Predicciones de entrenamiento
            oTrainPredict = nn_model.predict(train_x)
            fTrainingFScore = f1_score(train_y, oTrainPredict)
            fTrainingLogLoss = log_loss(train_y, oTrainPredict)

            # Predicciones de validacion
            oValidationPredict = nn_model.predict(oCleanValidationData.drop(columns = ['Survived']))
            fValidationFScore = f1_score(oCleanValidationData['Survived'], oValidationPredict)
            fValidationLogLoss = log_loss(oCleanValidationData['Survived'], oValidationPredict)

            oNNScore.append(fValidationFScore)

            if fValidationFScore > best_score:
                best_score = fValidationFScore
                best_nn_model = nn_model
                best_precision = precision_score(oCleanValidationData['Survived'], oValidationPredict)
                best_results.append({
                    "layers": layer,
                    "train_size": train_size,
                    "solver": solver,
                    "precision": best_precision,
                    "f1_score": best_score,
                    "log_loss": fValidationLogLoss,
                })

        results.append((layer, train_size, solver, fTrainingFScore, fTrainingLogLoss,
```



```
nn_model_data = pd.DataFrame(results, columns = ["layers", "train_size", "solver", "f1_score", "log_loss", "precision"])
nn_best_results_data = pd.DataFrame(best_results)
nn_best_results_data.to_csv("nn_best_results.csv", index=False)
print(nn_best_results_data)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:691: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (2000) reached and the optimization hasn't converged yet.
```

```
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:691: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (2000) reached and the optimization hasn't converged yet.
```

```
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:691: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (2000) reached and the optimization hasn't converged yet.
```

```
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:691: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (2000) reached and the optimization hasn't converged yet.
```


```
warnings.warn(
layers train_size solver precision f1_score log_loss
0 (2, 2, 2, 1) 0.2 adam 0.000000 0.000000 14.478209
1 (2, 2, 2, 1) 0.2 sgd 0.401685 0.573146 21.565444
2 (2, 2, 2, 1) 0.3 lbfgs 0.793103 0.710425 7.593466
3 (2, 2, 2, 1) 0.4 adam 0.726619 0.716312 8.099697
4 (2, 2, 2, 1) 0.5 lbfgs 0.739437 0.736842 7.593466
5 (2, 2, 2, 1) 0.8 adam 0.808333 0.737643 6.985989
6 (2, 2, 2, 1) 0.9 lbfgs 0.766917 0.739130 7.289728
7 (2, 3, 2, 1) 0.6 adam 0.766423 0.750000 7.087235
8 (3, 4, 3, 1) 0.6 adam 0.773723 0.757143 6.884743
9 (5, 4, 2, 1) 0.4 adam 0.829268 0.766917 6.277265
10 (6, 2, 2, 1) 0.6 lbfgs 0.848739 0.770992 6.074773
```

```
In [ ]: best_index = nn_best_results_data["f1_score"].idxmax()
best_nn_model = best_nn_model # Utiliza el modelo del mejor F1 score global
test_predicciones = best_nn_model.predict(oCleanTestData.drop(columns = ['Survived']))


fTestFScore = f1_score(oCleanTestData["Survived"], test_predicciones)
print("F1 Score: ", fTestFScore)
print("Log Loss: ", log_loss(oCleanTestData["Survived"], test_predicciones))
print("Precision: ", precision_score(oCleanTestData["Survived"], test_predicciones))

cm = confusion_matrix(oCleanTestData["Survived"], test_predicciones)
disp = ConfusionMatrixDisplay(confusion_matrix = cm, display_labels = best_model.class_labels)
disp.plot(cmap = plt.cm.Blues)
plt.show()
```

```
F1 Score: 0.7669172932330827
Log Loss: 6.242196955657161
Precision: 0.7727272727272727
```



```
In [ ]: plt.figure(figsize=(10, 6))
plt.scatter(nn_model_data["train_size"] , nn_model_data["fTrainingFScore"], label='Train')
plt.scatter(nn_model_data["train_size"] , nn_model_data["fValidationFScore"], label='Validation')
plt.axhline(y=fTestFScore, color='red', linestyle='--', label='Test Error (Base)')
plt.xlabel('Tamaño de muestra')
plt.ylabel('Error cuadrático medio')
plt.title('Evolución del error promedio de entrenamiento y validación')
plt.legend()
plt.show()
```



```
In [ ]: nn_model_data = nn_model_data.sort_values(by = "fValidationFScore", ascending = False)
nn_model_data.head()
```

Out []:

	layers	train_size	solver	fTrainingFScore	fTrainingLogLoss	fValidationFScore	fValidationLogLoss
590	(6, 2, 2, 1)	0.6	lbfgs	0.842912	5.277821	0.770992	6.074773
534	(5, 4, 2, 1)	0.4	adam	0.848485	5.119837	0.766917	6.277265
563	(5, 4, 3, 1)	0.5	lbfgs	0.906667	3.290942	0.762264	6.378512
642	(6, 3, 2, 1)	0.8	adam	0.809249	6.644919	0.758364	6.581004
690	(6, 4, 2, 1)	0.8	adam	0.835655	5.907154	0.757895	6.985989

```
In [ ]: nn_best_results_data = nn_best_results_data.sort_values(by = "f1_score", ascending = False)
nn_best_results_data.head()
```

Out[]:	layers	train_size	solver	precision	f1_score	log_loss
10	(6, 2, 2, 1)	0.6	lbfgs	0.848739	0.770992	6.074773
9	(5, 4, 2, 1)	0.4	adam	0.829268	0.766917	6.277265
8	(3, 4, 3, 1)	0.6	adam	0.773723	0.757143	6.884743
7	(2, 3, 2, 1)	0.6	adam	0.766423	0.750000	7.087235
6	(2, 2, 2, 1)	0.9	lbfgs	0.766917	0.739130	7.289728

```
In [ ]: oTestData = pd.read_csv('/content/test.csv')

oSubmissionTestData = clean_data(oTestData)
best_index = nn_best_results_data["f1_score"].idxmax()
best_nn_model = best_nn_model # Utiliza el modelo del mejor F1 score global
test_predicciones = best_nn_model.predict(oSubmissionTestData)

# Crear el archivo de salida con las predicciones
submit_data = pd.DataFrame(oTestData, columns=["PassengerId"])
submit_data["Survived"] = test_predicciones
submit_data.to_csv("prediction_nn.csv", index=False)
print("Predicciones guardadas en 'prediction.csv'")
```

Predicciones guardadas en 'prediction.csv'

Modelo Arbol de Decisión

```
In [ ]: import numpy as np
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import f1_score, log_loss, precision_score

# Inicialización
oFScore = []
train_sizes = np.arange(0.2, 1.0, 0.1) # 20% a 90%, incrementando de 10 en 10

best_results = []
best_model = None
best_max_depth = None
best_min_samples_leaf = None
max_depth_list = [3, 4, 5, 6, 7] # Profundidades máximas del árbol
min_samples_leaf_list = [1, 2, 4] # Número mínimo de muestras en una hoja
results = []

# Ciclo para probar diferentes hiperparámetros
for max_depth in max_depth_list:
    for min_samples_leaf in min_samples_leaf_list:
        for train_size in train_sizes:
            best_seed, best_score, best_precision = None, float("-inf"), float("-inf")
            best_score = 0

            # Generar 100 semillas aleatorias
            seeds = [[np.random.randint(0, 2**32 - 1), np.random.randint(0, 2**32 - 1)]
                    for seed in seeds:
                        # Tamaño del sample
```

```

iSampleSize = int(len(oCleanTrainData) * train_size)
oTrainSubset = oCleanTrainData.sample(n=iSampleSize)

# Oversampling para balancear las clases
x, y = oversample(oTrainSubset[["Age", "Pclass", "Sex", "SibSp", "Paro

# Entrenamiento del modelo utilizando DecisionTreeClassifier
oModelTemp = DecisionTreeClassifier(random_state=seed[1], max_depth=ma
oModelTemp.fit(x, y)

# Predicciones en el conjunto de entrenamiento
oTrainingPredict = oModelTemp.predict(x)
fTrainingFScore = f1_score(y, oTrainingPredict)
fTrainingLogLoss = log_loss(y, oTrainingPredict)

# Predicciones en el conjunto de validación
oValidationPredict = oModelTemp.predict(oCleanValidationData[["Age", "
fValidationFScore = f1_score(oCleanValidationData["Survived"], oValida
fValidationLogLoss = log_loss(oCleanValidationData["Survived"], oValic

# Almacenar el F-score de validación
oFScore.append(fValidationFScore)

# Actualizar el mejor modelo
if fValidationFScore > best_score:
    best_score = fValidationFScore
    best_seed = seed
    best_precision = precision_score(oCleanValidationData["Survived"],
    best_model = oModelTemp
    best_max_depth = max_depth
    best_min_samples_leaf = min_samples_leaf
    best_results.append({
        "train_size": train_size,
        "best_seed_split": best_seed[0],
        "best_seed_train": best_seed[1],
        "best_precision": best_precision,
        "best_score": best_score,
        "best_max_depth": best_max_depth,
        "best_min_samples_leaf": best_min_samples_leaf,
        "best_model": best_model,
    })

# Almacenar todos los resultados
results.append((train_size, max_depth, min_samples_leaf, fTrainingFSco

# Convertir resultados a DataFrame y exportar a CSV
decision_tree_data = pd.DataFrame(results, columns=["train_size", "max_depth", "min_sa
best_results_df = pd.DataFrame(best_results)
best_results_df.to_csv("best_tree_results.csv", index=False)

# Mostrar los mejores resultados
print(best_results_df)

```

	train_size	best_seed_split	best_seed_train	best_precision	best_score	\
0	0.2	334490763	3583404569	0.685484	0.646388	
1	0.2	1686234542	4089353493	0.641176	0.705502	
2	0.2	871169497	3583018481	0.713287	0.723404	
3	0.2	2152480510	934373311	0.668539	0.750789	
4	0.3	2465790219	3738454476	0.692308	0.669145	
..	
630	0.8	3198213124	1754648932	0.775194	0.746269	
631	0.9	3887606473	468304133	0.775862	0.705882	
632	0.9	2122526050	2637109213	0.757812	0.726592	
633	0.9	1465539036	372016571	0.755556	0.744526	
634	0.9	728753105	3385472702	0.770992	0.748148	

	best_max_depth	best_min_samples_leaf	\
0	3	1	
1	3	1	
2	3	1	
3	3	1	
4	3	1	
..	
630	7	4	
631	7	4	
632	7	4	
633	7	4	
634	7	4	

	best_model
0	DecisionTreeClassifier(max_depth=3, random_sta...
1	DecisionTreeClassifier(max_depth=3, random_sta...
2	DecisionTreeClassifier(max_depth=3, random_sta...
3	DecisionTreeClassifier(max_depth=3, random_sta...
4	DecisionTreeClassifier(max_depth=3, random_sta...
..	...
630	DecisionTreeClassifier(max_depth=7, min_sample...
631	DecisionTreeClassifier(max_depth=7, min_sample...
632	DecisionTreeClassifier(max_depth=7, min_sample...
633	DecisionTreeClassifier(max_depth=7, min_sample...
634	DecisionTreeClassifier(max_depth=7, min_sample...

[635 rows x 8 columns]


```
In [ ]: # Evaluar el mejor modelo en el conjunto de prueba
best_index = best_results_df["best_score"].idxmax()
best_model = best_results_df.iloc[best_index]["best_model"]

test_predicciones = best_model.predict(oCleanTestData[["Age", "Pclass", "Sex", "SibSp", "Parch", "Fare", "Embarked"]])

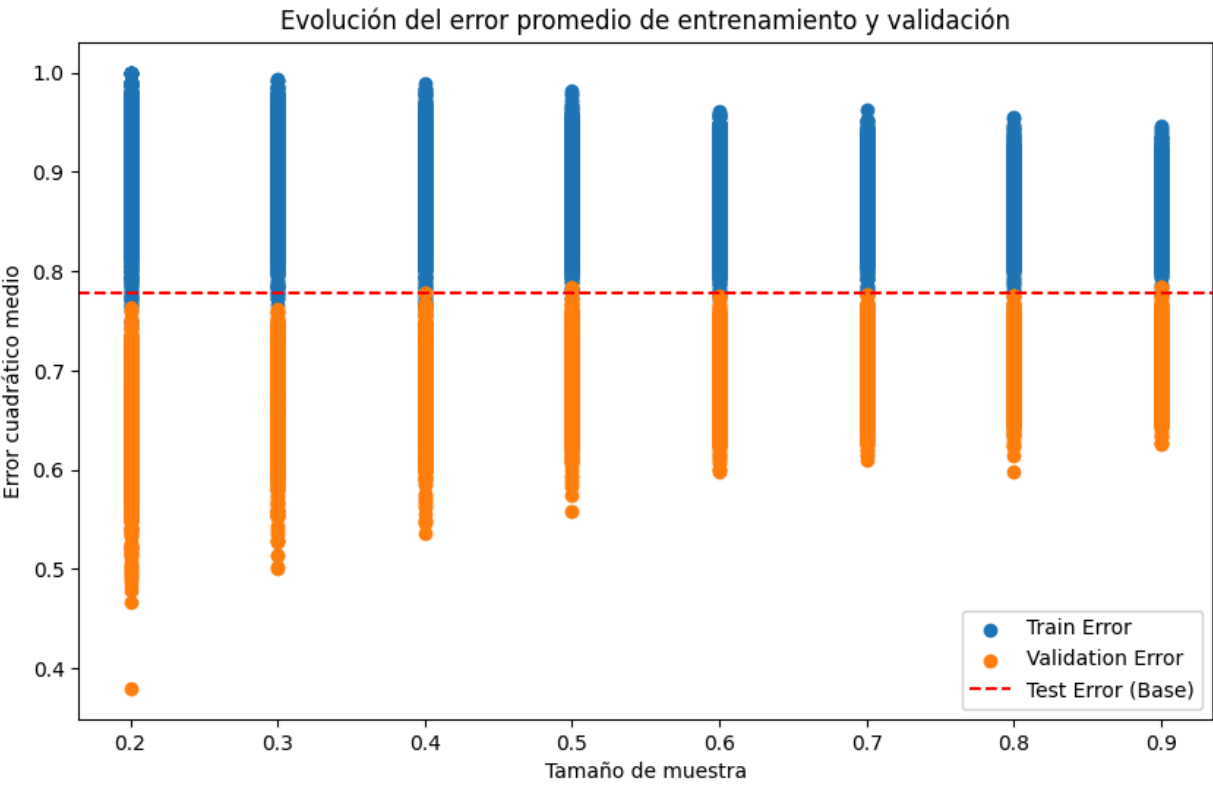
fTestFScore = f1_score(oCleanTestData["Survived"], test_predicciones)
print("F1 Score: ", fTestFScore)
print("Log Loss: ", log_loss(oCleanTestData["Survived"], test_predicciones))
print("Precision: ", precision_score(oCleanTestData["Survived"], test_predicciones))

cm = confusion_matrix(oCleanTestData["Survived"], test_predicciones)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=best_model.classes_)
disp.plot(cmap=plt.cm.Blues)
plt.show()

F1 Score:    0.7785234899328859
Log Loss:    6.6449193398931055
Precision:    0.7631578947368421
```



```
In [ ]: plt.figure(figsize=(10, 6))
plt.scatter(decision_tree_data["train_size"], decision_tree_data["fTrainingFScore"],
plt.scatter(decision_tree_data["train_size"], decision_tree_data["fValidationFScore"]
plt.axhline(y=fTestFScore, color='red', linestyle='--', label='Test Error (Base)')
plt.xlabel('Tamaño de muestra')
plt.ylabel('Error cuadrático medio')
plt.title('Evolución del error promedio de entrenamiento y validación')
plt.legend()
plt.show()
```



```
In [ ]: decision_tree_data = decision_tree_data.sort_values(by = "fValidationFScore", ascending=True)
decision_tree_data.head()
```

Out[]:

	train_size	max_depth	min_samples_leaf	fTrainingFScore	fTrainingLogLoss	fValidationFScore	f
6483	0.2	5	4	0.886364	3.834431	0.379888	
6494	0.2	5	4	0.888889	3.887061	0.466368	
2488	0.2	4	1	0.912621	3.243929	0.478431	
4005	0.2	4	4	0.839506	4.984761	0.483412	
8061	0.2	6	2	0.919540	2.867109	0.486692	

```
In [ ]: best_results_df = best_results_df.sort_values(by="best_score", ascending=False)
best_results_df.head()
```


Out[]:	train_size	best_seed_split	best_seed_train	best_precision	best_score	best_max_depth	best_min_
141	0.5	124441603	1850611572	0.770833	0.784452	4	
201	0.9	3025508625	3843109387	0.770833	0.784452	4	
241	0.9	1560008029	3270805070	0.738854	0.783784	4	
221	0.5	3837535486	2002760892	0.824000	0.780303	4	
183	0.5	792563679	332139667	0.824000	0.780303	4	

```
In [ ]: # Cargar los datos de prueba
oTestData = pd.read_csv('/content/test.csv')

# Limpiar los datos de prueba utilizando la función clean_data
oSubmissionTestData = clean_data(oTestData)

# Encontrar el mejor índice del modelo basado en el mejor F1 score
best_index = best_results_df["best_score"].idxmax()
best_model = best_results_df.iloc[best_index]["best_model"] # Recuperar el mejor modelo

# Realizar predicciones en los datos de prueba
test_predicciones = best_model.predict(oSubmissionTestData[["Age", "Pclass", "Sex", "SibSp", "Parch", "Fare", "Embarked"]])

# Crear el archivo de salida con las predicciones
submit_data = pd.DataFrame(oTestData["PassengerId"])
submit_data["Survived"] = test_predicciones
submit_data.to_csv("predictionAD.csv", index=False)

print("Predicciones guardadas en 'prediction.csv'")
```

Predicciones guardadas en 'prediction.csv'

Modelo Bosque Aleatorio

```
In [ ]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import f1_score, log_loss, precision_score
import numpy as np
import pandas as pd

# Definir las combinaciones de hiperparámetros
n_estimators_list = [100, 150, 200, 250, 300] # Número de árboles en el bosque
max_depth_list = [4, 5] # Máxima profundidad del árbol
train_sizes = np.arange(0.2, 1.0, 0.1) # Tamaños de entrenamiento del 20% al 90%

oFScore = []
results = []
best_results = []
best_model = None
best_params = None
```

```

# Iterar sobre los hiperparámetros
for n_estimators in n_estimators_list:
    for max_depth in max_depth_list:
        for train_size in train_sizes:
            print(f'n_estimators: {n_estimators}, max_depth: {max_depth}, size: {train_size}')

            best_seed, best_score, best_precision = None, float("-inf"), float("-inf")
            best_score = 0

            # Generar 100 pares de semillas para reproducibilidad
            seeds = [[np.random.randint(0, 2**32 - 1), np.random.randint(0, 2**32 - 1)] for _ in range(100)]
            for seed in seeds:

                # Asignar el tamaño del sample con el que se va a entrenar
                iSampleSize = int(len(oCleanTrainData) * train_size)

                # Separar el subset de entrenamiento según el tamaño asignado
                oTrainSubset = oCleanTrainData.sample(n=iSampleSize)

                x, y = oversample(oTrainSubset[["Age", "Pclass", "Sex", "SibSp", "Parch"]])

                # Entrenar el modelo RandomForest con los hiperparámetros
                oModelTemp = RandomForestClassifier(n_estimators=n_estimators, max_depth=max_depth)
                #oModelTemp.fit(oTrainSubset[["Age", "Pclass", "Sex", "SibSp", "Parch"]], y)
                oModelTemp.fit(x, y)

                # Generar predicciones con el modelo ya entrenado en el conjunto de entrenamiento
                oTrainingPredict = oModelTemp.predict(oTrainSubset[["Age", "Pclass", "Sex", "SibSp", "Parch"]])
                fTrainingFScore = f1_score(oTrainSubset["Survived"], oTrainingPredict)
                fTrainingLogLoss = log_loss(oTrainSubset["Survived"], oTrainingPredict)

                # Generar predicciones con el modelo en el conjunto de validación
                oValidationPredict = oModelTemp.predict(oCleanValidationData[["Age", "Pclass", "Sex", "SibSp", "Parch"]])
                fValidationFScore = f1_score(oCleanValidationData["Survived"], oValidationPredict)
                fValidationLogLoss = log_loss(oCleanValidationData["Survived"], oValidationPredict)

                # Guardar el F1-Score de validación
                oFScore.append(fValidationFScore)

                # Si el F1-Score de validación es mejor, guardar el mejor modelo y sus hiperparámetros
                if fValidationFScore > best_score:
                    best_score = fValidationFScore
                    best_seed = seed
                    best_precision = precision_score(oCleanValidationData["Survived"], oValidationPredict)
                    best_model = oModelTemp # Guardar el mejor modelo
                    best_params = {"n_estimators": n_estimators, "max_depth": max_depth}
                    best_results.append({
                        "train_size": train_size,
                        "best_seed_split": best_seed[0],
                        "best_seed_train": best_seed[1],
                        "best_precision": best_precision,
                        "best_score": best_score,
                        "n_estimators": n_estimators,
                        "max_depth": max_depth,
                    })

            # Guardar todos los resultados
            results.append((train_size, n_estimators, max_depth, fTrainingFScore, fTrainingLogLoss, best_score, best_precision, best_model, best_params, best_results))

```

```
# Convertir resultados a DataFrame y exportar a CSV
rf_data = pd.DataFrame(results, columns=["train_size", "n_estimators", "max_depth", "f
best_results_df = pd.DataFrame(best_results)
best_results_df.to_csv("best_rf_results.csv", index=False)

# Mostrar los mejores resultados
print(best_results_df)
```

36/40

```

n_estimators: 250, max_depth: 5, size: 0.6000000000000001
n_estimators: 250, max_depth: 5, size: 0.7000000000000002
n_estimators: 250, max_depth: 5, size: 0.8000000000000003
n_estimators: 250, max_depth: 5, size: 0.9000000000000001
n_estimators: 300, max_depth: 4, size: 0.2
n_estimators: 300, max_depth: 4, size: 0.3000000000000004
n_estimators: 300, max_depth: 4, size: 0.4000000000000001
n_estimators: 300, max_depth: 4, size: 0.5000000000000001
n_estimators: 300, max_depth: 4, size: 0.6000000000000001
n_estimators: 300, max_depth: 4, size: 0.7000000000000002
n_estimators: 300, max_depth: 4, size: 0.8000000000000003
n_estimators: 300, max_depth: 4, size: 0.9000000000000001
n_estimators: 300, max_depth: 5, size: 0.2
n_estimators: 300, max_depth: 5, size: 0.3000000000000004
n_estimators: 300, max_depth: 5, size: 0.4000000000000001
n_estimators: 300, max_depth: 5, size: 0.5000000000000001
n_estimators: 300, max_depth: 5, size: 0.6000000000000001
n_estimators: 300, max_depth: 5, size: 0.7000000000000002
n_estimators: 300, max_depth: 5, size: 0.8000000000000003
n_estimators: 300, max_depth: 5, size: 0.9000000000000001
  train_size  best_seed_split  best_seed_train  best_precision  best_score \
0           0.2        1049225254        351134135        0.726316    0.589744
1           0.2        1615704700        3401826984        0.786325    0.718750
2           0.2        2196881482        3961196112        0.923913    0.735931
3           0.2        1871638462        3961373903        0.741007    0.741007
4           0.2        2585428943        2079693359        0.733333    0.761246
..          ...          ...          ...          ...          ...
410         0.8        2603120915        890619019        0.830645    0.783270
411         0.9        3066529332        319564215        0.800000    0.773234
412         0.9        561756933        3129887841        0.827869    0.773946
413         0.9        413504734        193924108        0.818898    0.781955
414         0.9        2637130062        2991540583        0.832000    0.787879

  n_estimators  max_depth
0            100         4
1            100         4
2            100         4
3            100         4
4            100         4
..          ...          ...
410           300         5
411           300         5
412           300         5
413           300         5
414           300         5

```

[415 rows x 7 columns]

```

In [ ]: # Obtener el índice del mejor resultado (basado en el F1 Score más alto)
best_index = best_results_df["best_score"].idxmax()

# Utilizar el mejor modelo almacenado en 'best_model' para predecir en los datos de prueba
test_predicciones = best_model.predict(oCleanTestData[["Age", "Pclass", "Sex", "SibSp"]])

# Calcular las métricas de evaluación en los datos de prueba
fTestFScore = f1_score(oCleanTestData["Survived"], test_predicciones)
fTestLogLoss = log_loss(oCleanTestData["Survived"], test_predicciones)
fTestPrecision = precision_score(oCleanTestData["Survived"], test_predicciones)

# Mostrar las métricas de evaluación

```


```

print("F1 Score: ", fTestFScore)
print("Log Loss: ", fTestLogLoss)
print("Precision: ", fTestPrecision)

# Crear y mostrar la matriz de confusión
cm = confusion_matrix(oCleanTestData["Survived"], test_predicciones)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=best_model.classes_)
disp.plot(cmap=plt.cm.Blues)
plt.show()

```


F1 Score: 0.6614173228346456
 Log Loss: 8.658531261072836
 Precision: 0.6885245901639344



```

In [ ]: plt.figure(figsize=(10, 6))
plt.scatter(rf_data["train_size"], rf_data["fTrainingFScore"], label='Train Error')
plt.scatter(rf_data["train_size"], rf_data["fValidationFScore"], label='Validation Error')
plt.axhline(y=fTestFScore, color='red', linestyle='--', label='Test Error (Base)')
plt.xlabel('Tamaño de muestra')
plt.ylabel('Error cuadrático medio')
plt.title('Evolución del error promedio de entrenamiento y validación')
plt.legend()
plt.show()

```



```
In [ ]: rf_data = rf_data.sort_values(by = "fValidationFScore", ascending = True)
rf_data.head()
```

Out[]:

	train_size	n_estimators	max_depth	fTrainingFScore	fTrainingLogLoss	fValidationFScore	fValid
5602	0.2	250	5	0.888889	3.045943	0.502092	
29	0.2	100	4	0.742857	4.568914	0.538117	
5685	0.2	250	5	0.867925	3.553600	0.553571	
7357	0.3	300	5	0.840580	3.740379	0.579926	
0	0.2	100	4	0.818182	4.061257	0.589744	

```
In [ ]: best_results_df = best_results_df.sort_values(by = "best_score", ascending = True)
best_results_df.head()
```

Out[]:

	train_size	best_seed_split	best_seed_train	best_precision	best_score	best_max_depth	best_min_
159	0.2	1119926316	1047122078	0.719101	0.561404	4	
425	0.2	343772006	243201809	0.693878	0.573840	6	
333	0.2	3540181433	1634887774	0.701031	0.576271	5	
287	0.2	2161652909	4212499901	0.755814	0.577778	5	
242	0.2	428893739	2106457437	0.666667	0.600791	5	

In []:

```

import pandas as pd

# Leer el archivo de datos de prueba
oTestData = pd.read_csv('/content/test.csv')

# Limpiar los datos de prueba utilizando la función de limpieza que ya has definido
oSubmissionTestData = clean_data(oTestData)

# Utilizar el modelo del mejor F1 score global
best_index = best_results_df["best_score"].idxmax()
best_model = best_model # El modelo con el mejor F1 score global

# Hacer predicciones con el mejor modelo
test_predicciones = best_model.predict(oSubmissionTestData[["Age", "Pclass", "Sex", "S

# Crear el archivo de salida con las predicciones
submit_data = pd.DataFrame({
    "PassengerId": oTestData["PassengerId"],
    "Survived": test_predicciones
})
submit_data.to_csv("predictionRF.csv", index=False)

print("Predicciones guardadas en 'prediction.csv'")

```

Predicciones guardadas en 'prediction.csv'